

BLINK DB

1.0

Generated by Doxygen 1.13.2



---

<b>1 BLINK DB Documentation</b>	<b>1</b>
1.1 Overview . . . . .	1
1.2 Architecture . . . . .	1
1.2.1 Storage Engine (Part A) . . . . .	1
1.2.2 Advanced Storage Engine (Part B) . . . . .	2
1.2.3 Network Server (Part B) . . . . .	2
1.3 Performance Characteristics . . . . .	2
1.3.1 Benchmark Results . . . . .	2
1.3.2 Key Features . . . . .	2
1.4 Getting Started . . . . .	3
1.4.1 Building the Project . . . . .	3
1.4.2 Running the Server . . . . .	3
1.4.3 Running the Client . . . . .	3
1.4.4 Running Benchmarks . . . . .	3
1.4.5 Running Part A REPL . . . . .	3
1.5 Implementation Differences . . . . .	3
1.5.1 Part A vs Part B Storage Engines . . . . .	3
1.5.2 Key Classes in Part B . . . . .	3
1.6 Project Structure . . . . .	4
1.7 API Documentation . . . . .	4
1.7.1 Basic Storage Engine API (Part A) . . . . .	4
1.7.2 Advanced Storage Engine API (Part B) . . . . .	5
1.7.3 LRUCache API (Part B) . . . . .	5
1.7.4 Network Server API (Part B) . . . . .	5
1.7.5 Network Client API (Part B) . . . . .	6
1.8 Performance Optimization . . . . .	6
1.8.1 Memory Management (Part B) . . . . .	6
1.8.2 Disk Operations . . . . .	6
1.8.3 Network Optimization (Part B) . . . . .	6
1.9 Future Improvements . . . . .	6
1.9.1 Performance . . . . .	6
1.9.2 Features . . . . .	7
1.9.3 Reliability . . . . .	7
<b>2 Class Index</b>	<b>9</b>
2.1 Class List . . . . .	9
<b>3 File Index</b>	<b>11</b>
3.1 File List . . . . .	11
<b>4 Class Documentation</b>	<b>13</b>
4.1 StorageEngine::BatchEntry Struct Reference . . . . .	13
4.1.1 Member Data Documentation . . . . .	13

---

4.1.1.1 key . . . . .	13
4.1.1.2 value . . . . .	13
4.2 StorageEngine::DiskEntry Struct Reference . . . . .	13
4.2.1 Member Data Documentation . . . . .	14
4.2.1.1 offset . . . . .	14
4.2.1.2 size . . . . .	14
4.3 DiskStorage Class Reference . . . . .	14
4.3.1 Constructor & Destructor Documentation . . . . .	15
4.3.1.1 DiskStorage() . . . . .	15
4.3.1.2 ~DiskStorage() . . . . .	15
4.3.2 Member Function Documentation . . . . .	15
4.3.2.1 ensure_directory_exists() . . . . .	15
4.3.2.2 get() . . . . .	16
4.3.2.3 get_executable_path() . . . . .	16
4.3.2.4 load_data() . . . . .	16
4.3.2.5 put() . . . . .	17
4.3.2.6 remove() . . . . .	17
4.3.2.7 save_data() . . . . .	17
4.3.3 Member Data Documentation . . . . .	18
4.3.3.1 data_ . . . . .	18
4.3.3.2 data_file_ . . . . .	18
4.3.3.3 mutex_ . . . . .	18
4.4 LRUCache Class Reference . . . . .	18
4.4.1 Constructor & Destructor Documentation . . . . .	19
4.4.1.1 LRUCache() . . . . .	19
4.4.1.2 ~LRUCache() . . . . .	19
4.4.2 Member Function Documentation . . . . .	20
4.4.2.1 capacity() . . . . .	20
4.4.2.2 evict_lru() . . . . .	20
4.4.2.3 get() . . . . .	20
4.4.2.4 move_to_front() . . . . .	21
4.4.2.5 put() . . . . .	21
4.4.2.6 remove() . . . . .	21
4.4.2.7 size() . . . . .	21
4.4.3 Member Data Documentation . . . . .	22
4.4.3.1 cache_ . . . . .	22
4.4.3.2 capacity_ . . . . .	22
4.4.3.3 head_ . . . . .	22
4.4.3.4 mutex_ . . . . .	22
4.4.3.5 tail_ . . . . .	22
4.5 NetworkClient Class Reference . . . . .	22
4.5.1 Constructor & Destructor Documentation . . . . .	23

---

4.5.1.1 <code>NetworkClient()</code>	23
4.5.1.2 <code>~NetworkClient()</code>	23
4.5.2 Member Function Documentation	23
4.5.2.1 <code>parse_response()</code>	23
4.5.2.2 <code>send_command()</code>	23
4.5.3 Member Data Documentation	24
4.5.3.1 <code>server_addr</code>	24
4.5.3.2 <code>sock</code>	24
4.6 LRUCache::Node Struct Reference	24
4.6.1 Constructor & Destructor Documentation	25
4.6.1.1 <code>Node()</code>	25
4.6.2 Member Data Documentation	25
4.6.2.1 <code>key</code>	25
4.6.2.2 <code>next</code>	25
4.6.2.3 <code>prev</code>	25
4.6.2.4 <code>value</code>	25
4.7 Server Class Reference	25
4.7.1 Constructor & Destructor Documentation	26
4.7.1.1 <code>Server()</code>	26
4.7.1.2 <code>~Server()</code>	27
4.7.2 Member Function Documentation	27
4.7.2.1 <code>encode_resp()</code>	27
4.7.2.2 <code>handle_client_data()</code>	27
4.7.2.3 <code>handle_new_connection()</code>	28
4.7.2.4 <code>process_command()</code>	28
4.7.2.5 <code>run()</code>	29
4.7.2.6 <code>set_nonblocking()</code>	29
4.7.2.7 <code>setup_kqueue()</code>	29
4.7.2.8 <code>setup_server()</code>	30
4.7.2.9 <code>stop()</code>	30
4.7.3 Member Data Documentation	30
4.7.3.1 <code>client_buffers</code>	30
4.7.3.2 <code>kq</code>	31
4.7.3.3 <code>LISTEN_BACKLOG</code>	31
4.7.3.4 <code>MAX_EVENTS</code>	31
4.7.3.5 <code>PORT</code>	31
4.7.3.6 <code>server_fd</code>	31
4.7.3.7 <code>should_stop</code>	31
4.7.3.8 <code>storage</code>	31
4.8 StorageEngine Class Reference	31
4.8.1 Constructor & Destructor Documentation	33
4.8.1.1 <code>StorageEngine() [1/2]</code>	33

4.8.1.2 ~StorageEngine() [1/2] . . . . .	33
4.8.1.3 StorageEngine() [2/2] . . . . .	33
4.8.1.4 ~StorageEngine() [2/2] . . . . .	34
4.8.2 Member Function Documentation . . . . .	34
4.8.2.1 async_write_worker() . . . . .	34
4.8.2.2 clear() [1/2] . . . . .	34
4.8.2.3 clear() [2/2] . . . . .	34
4.8.2.4 del() [1/2] . . . . .	35
4.8.2.5 del() [2/2] . . . . .	35
4.8.2.6 flush_write_buffer() . . . . .	35
4.8.2.7 force_flush() [1/2] . . . . .	36
4.8.2.8 force_flush() [2/2] . . . . .	36
4.8.2.9 get() [1/3] . . . . .	37
4.8.2.10 get() [2/3] . . . . .	37
4.8.2.11 get() [3/3] . . . . .	37
4.8.2.12 load_disk_index() . . . . .	38
4.8.2.13 pending_write_count() . . . . .	38
4.8.2.14 put() . . . . .	38
4.8.2.15 remove_from_disk_index() . . . . .	39
4.8.2.16 save_disk_index() . . . . .	39
4.8.2.17 set() [1/2] . . . . .	40
4.8.2.18 set() [2/2] . . . . .	40
4.8.2.19 size() [1/2] . . . . .	41
4.8.2.20 size() [2/2] . . . . .	41
4.8.2.21 stop_async_writer() . . . . .	41
4.8.2.22 sync() . . . . .	41
4.8.2.23 update_disk_index() . . . . .	41
4.8.3 Member Data Documentation . . . . .	42
4.8.3.1 access_order . . . . .	42
4.8.3.2 BATCH_SIZE . . . . .	42
4.8.3.3 cache_ . . . . .	42
4.8.3.4 data_ . . . . .	42
4.8.3.5 DATA_FILE . . . . .	42
4.8.3.6 DISK_DIR . . . . .	43
4.8.3.7 disk_index . . . . .	43
4.8.3.8 disk_storage_ . . . . .	43
4.8.3.9 INDEX_FILE . . . . .	43
4.8.3.10 MAX_CACHE_SIZE . . . . .	43
4.8.3.11 MAX_KEY_SIZE . . . . .	43
4.8.3.12 MAX_VALUE_SIZE . . . . .	43
4.8.3.13 mutex_ . . . . .	43
4.8.3.14 pending_writes . . . . .	43

---

4.8.3.15 running_ . . . . .	43
4.8.3.16 write_buffer . . . . .	44
4.8.3.17 write_cv_ . . . . .	44
4.8.3.18 write_mutex_ . . . . .	44
4.8.3.19 write_queue_ . . . . .	44
4.8.3.20 write_thread_ . . . . .	44
<b>5 File Documentation</b>	<b>45</b>
5.1 docs/mainpage.md File Reference . . . . .	45
5.2 part-a/src/repl.cpp File Reference . . . . .	45
5.2.1 Function Documentation . . . . .	46
5.2.1.1 main() . . . . .	46
5.2.1.2 printUsage() . . . . .	46
5.3 part-a/src/storage_engine.cpp File Reference . . . . .	46
5.4 part-b/src/storage_engine.cpp File Reference . . . . .	47
5.5 part-a/src/storage_engine.h File Reference . . . . .	47
5.6 storage_engine.h . . . . .	48
5.7 part-b/src/storage_engine.h File Reference . . . . .	49
5.8 storage_engine.h . . . . .	50
5.9 part-b/src/main_server.cpp File Reference . . . . .	55
5.9.1 Function Documentation . . . . .	55
5.9.1.1 main() . . . . .	55
5.9.1.2 signal_handler() . . . . .	56
5.9.2 Variable Documentation . . . . .	56
5.9.2.1 g_server . . . . .	56
5.10 part-b/src/network_client.cpp File Reference . . . . .	56
5.10.1 Function Documentation . . . . .	57
5.10.1.1 main() . . . . .	57
5.11 part-b/src/network_server.cpp File Reference . . . . .	57
5.11.1 Macro Definition Documentation . . . . .	57
5.11.1.1 TCP_NODELAY . . . . .	57
5.12 part-b/src/network_server.h File Reference . . . . .	58
5.13 network_server.h . . . . .	58
<b>Index</b>	<b>61</b>



# Chapter 1

# BLINK DB Documentation

## 1.1 Overview

BLINK DB is a high-performance key-value store implementation with both in-memory and disk-based storage capabilities. The project is divided into two main parts:

### 1. Part A: Basic Storage Engine

- Simple in-memory storage with access order tracking
- Basic disk-based persistence with binary format
- Synchronous write operations
- Simple disk index management
- REPL (Read-Eval-Print Loop) interface

### 2. Part B: Advanced Network Server

- Redis protocol (RESP) support
- Non-blocking I/O using kqueue/epoll
- High-concurrency client handling
- Advanced LRU cache implementation
- Asynchronous disk operations
- Network client implementation

## 1.2 Architecture

### 1.2.1 Storage Engine (Part A)

The basic storage engine provides core functionality for storing and retrieving key-value pairs. It features:

- Simple in-memory hash map storage
- Access order tracking for basic LRU behavior
- Binary format disk persistence
- Synchronous write operations
- Simple disk index management
- REPL interface for interactive testing

### 1.2.2 Advanced Storage Engine (Part B)

The advanced storage engine includes sophisticated caching and async operations:

- **LRUCache class:** Thread-safe LRU cache with doubly-linked list implementation
- **DiskStorage class:** Persistent storage with automatic directory management
- **StorageEngine class:** Main engine with async write worker thread
- Write buffering and batch operations
- Cross-platform executable path detection

### 1.2.3 Network Server (Part B)

The network server implements the Redis protocol and provides:

- Non-blocking I/O using kqueue/epoll
- High-concurrency client handling
- RESP protocol implementation
- Signal handling for graceful shutdown
- Network client implementation for testing

## 1.3 Performance Characteristics

### 1.3.1 Benchmark Results

1. **Low Load (1000 requests, 10 connections)**
  - SET: 66,666 ops/sec
  - GET: 83,333 ops/sec
2. **High Load (100,000 requests, 100 connections)**
  - SET: 167,785 ops/sec
  - GET: 168,067 ops/sec

### 1.3.2 Key Features

- High throughput under load
- Low latency for most operations
- Good scaling with concurrency
- Efficient resource utilization

## 1.4 Getting Started

### 1.4.1 Building the Project

```
# Part A: Storage Engine
cd part-a
make clean
make

# Part B: Network Server
cd part-b
make clean
make
```

### 1.4.2 Running the Server

```
cd part-b
./blinkdb_server
```

### 1.4.3 Running the Client

```
cd part-b
./blinkdb_client
```

### 1.4.4 Running Benchmarks

```
# Using redis-benchmark
redis-benchmark -p 9001 -n 100000 -c 100

# Using built-in benchmark
./benchmark 100000 100
```

### 1.4.5 Running Part A REPL

```
cd part-a
./blinkdb
```

## 1.5 Implementation Differences

### 1.5.1 Part A vs Part B Storage Engines

Feature	Part A (Basic)	Part B (Advanced)
<b>Storage Format</b>	Binary format (data.dat, index.dat)	Text format (data.txt)
<b>Caching</b>	Simple access order tracking	Thread-safe LRU cache with doubly-linked list
<b>Write Operations</b>	Synchronous, immediate flush	Asynchronous with background worker thread
<b>Thread Safety</b>	Basic mutex protection	Advanced thread-safe design with condition variables
<b>Disk Management</b>	Manual directory creation	Automatic executable path detection
<b>Memory Management</b>	Simple hash map	Sophisticated cache with eviction policies
<b>Interface</b>	REPL command-line	Network server with RESP protocol

### 1.5.2 Key Classes in Part B

- **LRUCache**: Implements thread-safe LRU eviction with O(1) operations

- **DiskStorage**: Handles persistent storage with automatic path management
- **StorageEngine**: Main engine coordinating cache and disk operations
- **Server**: Network server with kqueue/epoll for high concurrency
- **NetworkClient**: Client implementation for testing and benchmarking

## 1.6 Project Structure

```

+-- part-a/           # Basic Storage Engine
|   +-- src/
|   |   +-- storage_engine.cpp    # Basic storage implementation
|   |   +-- storage_engine.h     # Basic storage header
|   |   +-- repl.cpp            # REPL interface
|   +-- disk_storage/         # Disk storage files
|   |   +-- data.dat
|   |   +-- index.dat
|   +-- blinkdb              # Compiled executable
+-- part-b/           # Advanced Network Server
|   +-- src/
|   |   +-- main_server.cpp      # Server main entry point
|   |   +-- network_server.cpp  # Network server implementation
|   |   +-- network_server.h    # Network server header
|   |   +-- network_client.cpp  # Network client implementation
|   |   +-- storage_engine.cpp  # Advanced storage implementation
|   |   +-- storage_engine.h    # Advanced storage header
|   |   +-- benchmark.cpp       # Performance benchmark tool
|   |   +-- disk_storage/        # Disk storage files
|   |   +-- data.txt
|   +-- blinkdb_server        # Compiled server executable
|   +-- blinkdb_client         # Compiled client executable
|   +-- benchmark              # Compiled benchmark executable
+-- docs/               # Documentation
  +-- mainpage.md
  +-- doxygen/             # Generated documentation
    +-- html/
    +-- latex/

```

## 1.7 API Documentation

### 1.7.1 Basic Storage Engine API (Part A)

The basic storage engine provides the following main interfaces:

- `set(key, value)`: Store a key-value pair
- `get(key)`: Retrieve a value by key
- `del(key)`: Delete a key-value pair
- `clear()`: Remove all key-value pairs
- `force_flush()`: Force flush write buffer to disk
- `size()`: Get total number of entries

### 1.7.2 Advanced Storage Engine API (Part B)

The advanced storage engine provides enhanced functionality:

- `set(key, value)`: Store a key-value pair (alias for put)
- `put(key, value)`: Store a key-value pair with async write
- `get(key)`: Retrieve a value by key
- `get(key, value)`: Retrieve a value by key (reference version)
- `del(key)`: Delete a key-value pair
- `clear()`: Remove all key-value pairs
- `force_flush()`: Force flush pending writes
- `sync()`: Synchronize all pending operations
- `size()`: Get cache size
- `pending_write_count()`: Get number of pending writes
- `stop_async_writer()`: Stop the async write worker

### 1.7.3 LRUCache API (Part B)

The LRU cache provides thread-safe caching:

- `get(key, value)`: Retrieve value from cache
- `put(key, value)`: Store value in cache
- `remove(key)`: Remove key from cache
- `capacity()`: Get cache capacity
- `size()`: Get current cache size

### 1.7.4 Network Server API (Part B)

The network server implements:

- RESP protocol parsing and encoding
- Client connection management with kqueue/epoll
- Command processing (SET, GET, DEL, PING, etc.)
- Response formatting
- Signal handling for graceful shutdown

### 1.7.5 Network Client API (Part B)

The network client provides:

- TCP connection management
- RESP protocol communication
- Command sending and response parsing
- Error handling and connection testing

## 1.8 Performance Optimization

### 1.8.1 Memory Management (Part B)

- **LRU Cache:** Thread-safe doubly-linked list implementation for efficient cache management
- **Write Buffering:** Batch operations for better disk I/O performance
- **Memory-efficient Data Structures:** Optimized hash maps and linked lists
- **Async Write Worker:** Background thread for non-blocking disk operations

### 1.8.2 Disk Operations

- **Part A:** Synchronous binary format with simple index management
- **Part B:** Asynchronous writes with background worker thread
- **Cross-platform Path Detection:** Automatic executable path detection for storage location
- **Batch Index Updates:** Efficient disk index management

### 1.8.3 Network Optimization (Part B)

- **Non-blocking I/O:** kqueue/epoll for high-concurrency handling
- **RESP Protocol:** Efficient Redis-compatible protocol implementation
- **Signal Handling:** Graceful shutdown with SIGINT/SIGTERM support
- **Client Connection Management:** Efficient client buffer management

## 1.9 Future Improvements

### 1.9.1 Performance

- Implement connection pooling
- Add read-ahead caching
- Optimize disk index structure

### **1.9.2 Features**

- Add transaction support
- Implement pub/sub
- Add data compression

### **1.9.3 Reliability**

- Add data replication
- Implement checkpointing
- Add data validation



# Chapter 2

## Class Index

### 2.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

StorageEngine::BatchEntry . . . . .	13
StorageEngine::DiskEntry . . . . .	13
DiskStorage . . . . .	14
LRUCache . . . . .	18
NetworkClient . . . . .	22
LRUCache::Node . . . . .	24
Server . . . . .	25
StorageEngine . . . . .	31



# Chapter 3

## File Index

### 3.1 File List

Here is a list of all files with brief descriptions:

part-a/src/repl.cpp	45
part-a/src/storage_engine.cpp	46
part-a/src/storage_engine.h	47
part-b/src/main_server.cpp	55
part-b/src/network_client.cpp	56
part-b/src/network_server.cpp	57
part-b/src/network_server.h	58
part-b/src/storage_engine.cpp	47
part-b/src/storage_engine.h	49



# Chapter 4

## Class Documentation

### 4.1 StorageEngine::BatchEntry Struct Reference

#### Public Attributes

- std::string [key](#)
- std::string [value](#)

#### 4.1.1 Member Data Documentation

##### 4.1.1.1 [key](#)

```
std::string StorageEngine::BatchEntry::key
```

##### 4.1.1.2 [value](#)

```
std::string StorageEngine::BatchEntry::value
```

The documentation for this struct was generated from the following file:

- part-a/src/[storage\\_engine.h](#)

### 4.2 StorageEngine::DiskEntry Struct Reference

#### Public Attributes

- size\_t [offset](#)
- size\_t [size](#)

## 4.2.1 Member Data Documentation

### 4.2.1.1 offset

```
size_t StorageEngine::DiskEntry::offset
```

### 4.2.1.2 size

```
size_t StorageEngine::DiskEntry::size
```

The documentation for this struct was generated from the following file:

- part-a/src/storage\_engine.h

## 4.3 DiskStorage Class Reference

```
#include <storage_engine.h>
```

### Public Member Functions

- [DiskStorage \(\)](#)
- [~DiskStorage \(\)](#)
- [bool get \(const std::string &key, std::string &value\)](#)
- [bool put \(const std::string &key, const std::string &value\)](#)
- [void remove \(const std::string &key\)](#)

### Private Member Functions

- [std::filesystem::path get\\_executable\\_path \(\)](#)
- [void ensure\\_directory\\_exists \(\)](#)
- [void load\\_data \(\)](#)
- [void save\\_data \(\)](#)

### Private Attributes


- [std::string data\\_file\\_](#)
- [std::mutex mutex\\_](#)
- [std::unordered\\_map< std::string, std::string > data\\_](#)

### 4.3.1 Constructor & Destructor Documentation

#### 4.3.1.1 DiskStorage()

```
DiskStorage::DiskStorage () [inline]
```

Here is the call graph for this function:



#### 4.3.1.2 ~DiskStorage()

```
DiskStorage::~DiskStorage () [inline]
```

Here is the call graph for this function:



### 4.3.2 Member Function Documentation


#### 4.3.2.1 ensure\_directory\_exists()

```
void DiskStorage::ensure_directory_exists () [inline], [private]
```

Here is the call graph for this function:



Here is the caller graph for this function:




#### 4.3.2.2 `get()`

```
bool DiskStorage::get (
    const std::string & key,
    std::string & value) [inline]
```

#### 4.3.2.3 `get_executable_path()`

```
std::filesystem::path DiskStorage::get_executable_path () [inline], [private]
```

Here is the caller graph for this function:



#### 4.3.2.4 `load_data()`

```
void DiskStorage::load_data () [inline], [private]
```

Here is the caller graph for this function:



#### 4.3.2.5 put()

```
bool DiskStorage::put (
    const std::string & key,
    const std::string & value) [inline]
```

Here is the call graph for this function:



#### 4.3.2.6 remove()

```
void DiskStorage::remove (
    const std::string & key) [inline]
```


Here is the call graph for this function:



#### 4.3.2.7 save\_data()

```
void DiskStorage::save_data () [inline], [private]
```

Here is the caller graph for this function:



### 4.3.3 Member Data Documentation

#### 4.3.3.1 `data_`

```
std::unordered_map<std::string, std::string> DiskStorage::data_ [private]
```

#### 4.3.3.2 `data_file_`

```
std::string DiskStorage::data_file_ [private]
```

#### 4.3.3.3 `mutex_`

```
std::mutex DiskStorage::mutex_ [private]
```


The documentation for this class was generated from the following file:

- part-b/src/[storage\\_engine.h](#)

## 4.4 LRUCache Class Reference

```
#include <storage_engine.h>
```

Collaboration diagram for LRUCache:



## Classes

- struct [Node](#)

## Public Member Functions

- `LRUCache (size_t capacity)`
- `~LRUCache ()`
- `size_t capacity () const`
- `size_t size () const`
- `bool get (const std::string &key, std::string &value)`
- `bool put (const std::string &key, const std::string &value)`
- `bool remove (const std::string &key)`

## Private Member Functions

- `void move_to_front (Node *node)`
- `void evict_lru ()`

## Private Attributes

- `size_t capacity_`
- `std::unordered_map< std::string, Node * > cache_`
- `Node * head_`
- `Node * tail_`
- `std::mutex mutex_`

### 4.4.1 Constructor & Destructor Documentation

#### 4.4.1.1 LRUCache()

```
LRUCache::LRUCache (
    size_t capacity) [inline]
```

Here is the call graph for this function:



#### 4.4.1.2 ~LRUCache()

```
LRUCache::~LRUCache () [inline]
```

## 4.4.2 Member Function Documentation

### 4.4.2.1 capacity()

```
size_t LRUCache::capacity () const [inline]
```

Here is the caller graph for this function:



### 4.4.2.2 evict\_lru()

```
void LRUCache::evict_lru () [inline], [private]
```

Here is the caller graph for this function:



### 4.4.2.3 get()

```
bool LRUCache::get (
    const std::string & key,
    std::string & value) [inline]
```


Here is the call graph for this function:



#### 4.4.2.4 move\_to\_front()

```
void LRUCache::move_to_front (
    Node * node) [inline], [private]
```


Here is the caller graph for this function:



#### 4.4.2.5 put()

```
bool LRUCache::put (
    const std::string & key,
    const std::string & value) [inline]
```

Here is the call graph for this function:



#### 4.4.2.6 remove()

```
bool LRUCache::remove (
    const std::string & key) [inline]
```

#### 4.4.2.7 size()

```
size_t LRUCache::size () const [inline]
```

### 4.4.3 Member Data Documentation

#### 4.4.3.1 cache\_

```
std::unordered_map<std::string, Node\*> LRUCache::cache_ [private]
```

#### 4.4.3.2 capacity\_

```
size_t LRUCache::capacity_ [private]
```

#### 4.4.3.3 head\_

```
Node\* LRUCache::head_ [private]
```

#### 4.4.3.4 mutex\_

```
std::mutex LRUCache::mutex_ [private]
```

#### 4.4.3.5 tail\_

```
Node\* LRUCache::tail_ [private]
```

The documentation for this class was generated from the following file:

- part-b/src/[storage\\_engine.h](#)

## 4.5 NetworkClient Class Reference

### Public Member Functions

- [NetworkClient](#) (const char \*host="127.0.0.1", int port=9001)
- [~NetworkClient](#) ()
- std::string [send\\_command](#) (const std::string &command)

### Private Member Functions

- std::string [parse\\_response](#) (const std::string &resp)

### Private Attributes

- int [sock](#)
- struct sockaddr\_in [server\\_addr](#)

### 4.5.1 Constructor & Destructor Documentation

#### 4.5.1.1 NetworkClient()

```
NetworkClient::NetworkClient (
    const char * host = "127.0.0.1",
    int port = 9001) [inline]
```

#### 4.5.1.2 ~NetworkClient()

```
NetworkClient::~NetworkClient () [inline]
```

### 4.5.2 Member Function Documentation

#### 4.5.2.1 parse\_response()

```
std::string NetworkClient::parse_response (
    const std::string & resp) [inline], [private]
```

Here is the caller graph for this function:



#### 4.5.2.2 send\_command()

```
std::string NetworkClient::send_command (
    const std::string & command) [inline]
```

Here is the call graph for this function:



Here is the caller graph for this function:



### 4.5.3 Member Data Documentation

#### 4.5.3.1 server\_addr

```
struct sockaddr_in NetworkClient::server_addr [private]
```

#### 4.5.3.2 sock

```
int NetworkClient::sock [private]
```

The documentation for this class was generated from the following file:

- part-b/src/[network\\_client.cpp](#)

## 4.6 LRUCache::Node Struct Reference

Collaboration diagram for LRUCache::Node:



### Public Member Functions

- [Node](#) (const std::string &k, const std::string &v)

## Public Attributes

- std::string `key`
- std::string `value`
- `Node * prev`
- `Node * next`

### 4.6.1 Constructor & Destructor Documentation

#### 4.6.1.1 `Node()`

```
LRUCache::Node::Node (
    const std::string & k,
    const std::string & v) [inline]
```

### 4.6.2 Member Data Documentation

#### 4.6.2.1 `key`

```
std::string LRUCache::Node::key
```

#### 4.6.2.2 `next`

```
Node* LRUCache::Node::next
```

#### 4.6.2.3 `prev`

```
Node* LRUCache::Node::prev
```

#### 4.6.2.4 `value`

```
std::string LRUCache::Node::value
```

The documentation for this struct was generated from the following file:

- part-b/src/[storage\\_engine.h](#)

## 4.7 Server Class Reference

```
#include <network_server.h>
```

## Public Member Functions

- `Server ()`
- `~Server ()`
- `void run ()`
- `void stop ()`

## Private Member Functions

- `void setup_server ()`
- `void setup_kqueue ()`
- `void set_nonblocking (int fd)`
- `void handle_new_connection ()`
- `void handle_client_data (int client_fd)`
- `std::string process_command (const std::string &command)`
- `std::string encode_resp (const std::string &response)`

## Private Attributes

- `int server_fd = -1`
- `int kq = -1`
- `bool should_stop = false`
- `std::unique_ptr< StorageEngine > storage`
- `std::unordered_map< int, std::string > client_buffers`

## Static Private Attributes


- `static constexpr int PORT = 9001`
- `static constexpr int LISTEN_BACKLOG = 128`
- `static constexpr int MAX_EVENTS = 1024`

## 4.7.1 Constructor & Destructor Documentation

### 4.7.1.1 Server()

```
Server::Server ()
```

Here is the call graph for this function:



#### 4.7.1.2 ~Server()

```
Server::~Server ()
```

Here is the call graph for this function:



## 4.7.2 Member Function Documentation

#### 4.7.2.1 encode\_resp()

```
std::string Server::encode_resp (
    const std::string & response) [private]
```


#### 4.7.2.2 handle\_client\_data()

```
void Server::handle_client_data (
    int client_fd) [private]
```

Here is the call graph for this function:




Here is the caller graph for this function:



#### 4.7.2.3 handle\_new\_connection()

```
void Server::handle_new_connection () [private]
```

Here is the call graph for this function:




Here is the caller graph for this function:



#### 4.7.2.4 process\_command()

```
std::string Server::process_command (
    const std::string & command) [private]
```


Here is the caller graph for this function:



#### 4.7.2.5 run()

```
void Server::run ()
```


Here is the call graph for this function:



#### 4.7.2.6 set\_nonblocking()

```
void Server::set_nonblocking (
    int fd) [private]
```

Here is the caller graph for this function:



#### 4.7.2.7 setup\_kqueue()

```
void Server::setup_kqueue () [private]
```

Here is the caller graph for this function:



#### 4.7.2.8 `setup_server()`

```
void Server::setup_server () [private]
```

Here is the call graph for this function:




Here is the caller graph for this function:



#### 4.7.2.9 `stop()`

```
void Server::stop ()
```

Here is the caller graph for this function:



### 4.7.3 Member Data Documentation

#### 4.7.3.1 `client_buffers`

```
std::unordered_map<int, std::string> Server::client_buffers [private]
```

#### 4.7.3.2 kq

```
int Server::kq = -1 [private]
```

#### 4.7.3.3 LISTEN\_BACKLOG

```
int Server::LISTEN_BACKLOG = 128 [static], [constexpr], [private]
```

#### 4.7.3.4 MAX\_EVENTS

```
int Server::MAX_EVENTS = 1024 [static], [constexpr], [private]
```

#### 4.7.3.5 PORT

```
int Server::PORT = 9001 [static], [constexpr], [private]
```

#### 4.7.3.6 server\_fd

```
int Server::server_fd = -1 [private]
```

#### 4.7.3.7 should\_stop

```
bool Server::should_stop = false [private]
```

#### 4.7.3.8 storage

```
std::unique_ptr<StorageEngine> Server::storage [private]
```

The documentation for this class was generated from the following files:

- part-b/src/network\_server.h
- part-b/src/network\_server.cpp

## 4.8 StorageEngine Class Reference

```
#include <storage_engine.h>
```

### Classes

- struct [BatchEntry](#)
- struct [DiskEntry](#)

## Public Member Functions

- `StorageEngine ()`
- `~StorageEngine ()`
- `bool set (const std::string &key, const std::string &value)`
- `std::string get (const std::string &key)`
- `bool del (const std::string &key)`
- `void clear ()`
- `void force_flush ()`
- `size_t size () const`
- `StorageEngine (size_t cache_size=1)`
- `~StorageEngine ()`
- `bool set (const std::string &key, const std::string &value)`
- `std::string get (const std::string &key)`
- `bool get (const std::string &key, std::string &value)`
- `bool del (const std::string &key)`
- `void clear ()`
- `void force_flush ()`
- `size_t size () const`
- `void sync ()`
- `size_t pending_write_count () const`
- `void stop_async_writer ()`
- `bool put (const std::string &key, const std::string &value)`

## Private Member Functions

- `void load_disk_index ()`
- `void save_disk_index ()`
- `void update_disk_index (const std::string &key, size_t offset, size_t size)`
- `void remove_from_disk_index (const std::string &key)`
- `void flush_write_buffer ()`
- `void async_write_worker ()`

## Private Attributes

- `std::unordered_map< std::string, std::string > data_`
- `std::list< std::string > access_order`
- `std::mutex mutex_`
- `size_t pending_writes = 0`
- `std::map< std::string, DiskEntry > disk_index`
- `std::vector< BatchEntry > write_buffer`
- `std::unique_ptr< LRUcache > cache_`
- `std::unique_ptr< DiskStorage > disk_storage_`
- `std::queue< std::pair< std::string, std::string > > write_queue_`
- `std::mutex write_mutex_`
- `std::condition_variable write_cv_`
- `std::thread write_thread_`
- `std::atomic< bool > running_`

### Static Private Attributes

- static constexpr size\_t **MAX\_KEY\_SIZE** = 256
- static constexpr size\_t **MAX\_VALUE\_SIZE** = 1024
- static constexpr size\_t **MAX\_CACHE\_SIZE** = 10000000
- static constexpr size\_t **BATCH\_SIZE** = 1000000
- static constexpr const char \* **DISK\_DIR** = "disk\_storage"
- static constexpr const char \* **DATA\_FILE** = "data.dat"
- static constexpr const char \* **INDEX\_FILE** = "index.dat"

## 4.8.1 Constructor & Destructor Documentation

### 4.8.1.1 StorageEngine() [1/2]

`StorageEngine::StorageEngine ()`

Here is the call graph for this function:



### 4.8.1.2 ~StorageEngine() [1/2]

`StorageEngine::~StorageEngine ()`

Here is the call graph for this function:



### 4.8.1.3 StorageEngine() [2/2]

`StorageEngine::StorageEngine (
 size_t cache_size = 1) [inline]`

Here is the call graph for this function:



#### 4.8.1.4 ~StorageEngine() [2/2]

```
StorageEngine::~StorageEngine () [inline]
```

## 4.8.2 Member Function Documentation

### 4.8.2.1 async\_write\_worker()

```
void StorageEngine::async_write_worker () [inline], [private]
```

Here is the caller graph for this function:



### 4.8.2.2 clear() [1/2]

```
void StorageEngine::clear ()
```

Here is the caller graph for this function:



### 4.8.2.3 clear() [2/2]

```
void StorageEngine::clear () [inline]
```

#### 4.8.2.4 del() [1/2]

```
bool StorageEngine::del (
    const std::string & key)
```

Here is the call graph for this function:



Here is the caller graph for this function:




#### 4.8.2.5 del() [2/2]

```
bool StorageEngine::del (
    const std::string & key) [inline]
```


#### 4.8.2.6 flush\_write\_buffer()

```
void StorageEngine::flush_write_buffer () [private]
```

Here is the call graph for this function:




Here is the caller graph for this function:




#### 4.8.2.7 `force_flush()` [1/2]

```
void StorageEngine::force_flush ()
```

Here is the call graph for this function:



Here is the caller graph for this function:



#### 4.8.2.8 `force_flush()` [2/2]

```
void StorageEngine::force_flush () [inline]
```

**4.8.2.9 get() [1/3]**

```
std::string StorageEngine::get (
    const std::string & key)
```

Here is the caller graph for this function:

**4.8.2.10 get() [2/3]**

```
std::string StorageEngine::get (
    const std::string & key) [inline]
```

Here is the call graph for this function:


**4.8.2.11 get() [3/3]**

```
bool StorageEngine::get (
    const std::string & key,
    std::string & value) [inline]
```

#### 4.8.2.12 load\_disk\_index()

```
void StorageEngine::load_disk_index () [private]
```

Here is the call graph for this function:



Here is the caller graph for this function:



#### 4.8.2.13 pending\_write\_count()

```
size_t StorageEngine::pending_write_count () const [inline]
```

#### 4.8.2.14 put()

```
bool StorageEngine::put (
    const std::string & key,
    const std::string & value) [inline]
```

Here is the caller graph for this function:



#### 4.8.2.15 remove\_from\_disk\_index()

```
void StorageEngine::remove_from_disk_index (
    const std::string & key) [private]
```

Here is the call graph for this function:




Here is the caller graph for this function:



#### 4.8.2.16 save\_disk\_index()

```
void StorageEngine::save_disk_index () [private]
```

Here is the caller graph for this function:



#### 4.8.2.17 `set()` [1/2]

```
bool StorageEngine::set (
    const std::string & key,
    const std::string & value)
```

Here is the call graph for this function:



Here is the caller graph for this function:



#### 4.8.2.18 `set()` [2/2]

```
bool StorageEngine::set (
    const std::string & key,
    const std::string & value) [inline]
```


Here is the call graph for this function:



#### 4.8.2.19 size() [1/2]

```
size_t StorageEngine::size () const
```

Here is the caller graph for this function:



#### 4.8.2.20 size() [2/2]

```
size_t StorageEngine::size () const [inline]
```


#### 4.8.2.21 stop\_async\_writer()

```
void StorageEngine::stop_async_writer () [inline]
```

#### 4.8.2.22 sync()

```
void StorageEngine::sync () [inline]
```


Here is the call graph for this function:




#### 4.8.2.23 update\_disk\_index()

```
void StorageEngine::update_disk_index (
    const std::string & key,
    size_t offset,
    size_t size) [private]
```

Here is the call graph for this function:



Here is the caller graph for this function:



### 4.8.3 Member Data Documentation

#### 4.8.3.1 access\_order

```
std::list<std::string> StorageEngine::access_order [private]
```

#### 4.8.3.2 BATCH\_SIZE

```
size_t StorageEngine::BATCH_SIZE = 1000000 [static], [constexpr], [private]
```

#### 4.8.3.3 cache\_

```
std::unique_ptr<LRUCache> StorageEngine::cache_ [private]
```

#### 4.8.3.4 data\_

```
std::unordered_map<std::string, std::string> StorageEngine::data_ [private]
```

#### 4.8.3.5 DATA\_FILE

```
const char* StorageEngine::DATA_FILE = "data.dat" [static], [constexpr], [private]
```

#### 4.8.3.6 DISK\_DIR

```
const char* StorageEngine::DISK_DIR = "disk_storage" [static], [constexpr], [private]
```

#### 4.8.3.7 disk\_index

```
std::map<std::string, DiskEntry> StorageEngine::disk_index [private]
```

#### 4.8.3.8 disk\_storage\_

```
std::unique_ptr<DiskStorage> StorageEngine::disk_storage_ [private]
```

#### 4.8.3.9 INDEX\_FILE

```
const char* StorageEngine::INDEX_FILE = "index.dat" [static], [constexpr], [private]
```

#### 4.8.3.10 MAX\_CACHE\_SIZE

```
size_t StorageEngine::MAX_CACHE_SIZE = 10000000 [static], [constexpr], [private]
```

#### 4.8.3.11 MAX\_KEY\_SIZE

```
size_t StorageEngine::MAX_KEY_SIZE = 256 [static], [constexpr], [private]
```

#### 4.8.3.12 MAX\_VALUE\_SIZE

```
size_t StorageEngine::MAX_VALUE_SIZE = 1024 [static], [constexpr], [private]
```

#### 4.8.3.13 mutex\_

```
std::mutex StorageEngine::mutex_ [mutable], [private]
```

#### 4.8.3.14 pending\_writes

```
size_t StorageEngine::pending_writes = 0 [private]
```

#### 4.8.3.15 running\_

```
std::atomic<bool> StorageEngine::running_ [private]
```

#### 4.8.3.16 `write_buffer`

```
std::vector<BatchEntry> StorageEngine::write_buffer [private]
```

#### 4.8.3.17 `write_cv_`

```
std::condition_variable StorageEngine::write_cv_ [private]
```

#### 4.8.3.18 `write_mutex_`

```
std::mutex StorageEngine::write_mutex_ [private]
```

#### 4.8.3.19 `write_queue_`

```
std::queue<std::pair<std::string, std::string>> StorageEngine::write_queue_ [private]
```

#### 4.8.3.20 `write_thread_`

```
std::thread StorageEngine::write_thread_ [private]
```

The documentation for this class was generated from the following files:

- part-a/src/[storage\\_engine.h](#)
- part-b/src/[storage\\_engine.h](#)
- part-a/src/[storage\\_engine.cpp](#)


# Chapter 5

## File Documentation

### 5.1 docs/mainpage.md File Reference

### 5.2 part-a/src/repl.cpp File Reference

```
#include "storage_engine.h"
#include <iostream>
#include <sstream>
#include <string>
Include dependency graph for repl.cpp:
```



#### Functions


- void `printUsage ()`
- int `main ()`

## 5.2.1 Function Documentation

### 5.2.1.1 main()

```
int main ()
```

Here is the call graph for this function:



### 5.2.1.2 printUsage()

```
void printUsage ()
```


Here is the caller graph for this function:



## 5.3 part-a/src/storage\_engine.cpp File Reference


```
#include "storage_engine.h"
#include <fstream>
#include <iostream>
#include <filesystem>
#include <vector>
```

```
#include <sstream>
Include dependency graph for storage_engine.cpp:
```




## 5.4 part-b/src/storage\_engine.cpp File Reference

```
#include "storage_engine.h"
#include <iostream>
Include dependency graph for storage_engine.cpp:
```




## 5.5 part-a/src/storage\_engine.h File Reference

```
#include <string>
#include <unordered_map>
#include <mutex>
#include <list>
#include <iostream>
#include <map>
#include <vector>
Include dependency graph for storage_engine.h:
```



This graph shows which files directly or indirectly include this file:



## Classes

- class [StorageEngine](#)
- struct [StorageEngine::DiskEntry](#)
- struct [StorageEngine::BatchEntry](#)

## 5.6 storage\_engine.h

[Go to the documentation of this file.](#)

```

00001 #pragma once
00002
00003 #include <string>
00004 #include <unordered_map>
00005 #include <mutex>
00006 #include <list>
00007 #include <iostream>
00008 #include <map>
00009 #include <vector>
00010
00011 class StorageEngine {
00012 public:
00013     StorageEngine();
00014     ~StorageEngine();
00015
00016     bool set(const std::string& key, const std::string& value);
00017     std::string get(const std::string& key);
00018     bool del(const std::string& key);
00019     void clear(); // Clear all data from memory and disk
00020     void force_flush(); // Force flush write buffer
00021     size_t size() const; // Get total number of entries (in memory + on disk)
00022
00023 private:
00024     static constexpr size_t MAX_KEY_SIZE = 256;
00025     static constexpr size_t MAX_VALUE_SIZE = 1024;
00026     static constexpr size_t MAX_CACHE_SIZE = 10000000; // 10M entries max in memory
00027     static constexpr size_t BATCH_SIZE = 1000000; // 1M entries to batch write
00028     static constexpr const char* DISK_DIR = "disk_storage";
00029     static constexpr const char* DATA_FILE = "data.dat";
00030     static constexpr const char* INDEX_FILE = "index.dat";
00031
00032     struct DiskEntry {
00033         size_t offset;
00034         size_t size;
00035     };
00036
00037     struct BatchEntry {
00038         std::string key;
00039         std::string value;
  
```

```
00040      };
00041
00042      std::unordered_map<std::string, std::string> data_;
00043      std::list<std::string> access_order; // Track access order for LRU eviction
00044      mutable std::mutex mutex_; // Make mutex mutable for const methods
00045      size_t pending_writes = 0; // Track number of pending writes
00046      std::map<std::string, DiskEntry> disk_index; // Index for disk entries
00047      std::vector<BatchEntry> write_buffer; // Buffer for batch writes
00048
00049      void load_disk_index();
00050      void save_disk_index();
00051      void update_disk_index(const std::string& key, size_t offset, size_t size);
00052      void remove_from_disk_index(const std::string& key);
00053      void flush_write_buffer();
00054  };
```


## 5.7 part-b/src/storage\_engine.h File Reference

```
#include <string>
#include <unordered_map>
#include <mutex>
#include <list>
#include <iostream>
#include <map>
#include <vector>
#include <memory>
#include <thread>
#include <condition_variable>
#include <atomic>
#include <fstream>
#include <filesystem>
#include <queue>
```

Include dependency graph for storage\_engine.h:



This graph shows which files directly or indirectly include this file:



## Classes

- class [LRUCache](#)
- struct [LRUCache::Node](#)
- class [DiskStorage](#)
- class [StorageEngine](#)

## 5.8 storage\_engine.h

[Go to the documentation of this file.](#)

```

00001 #pragma once
00002
00003 #include <string>
00004 #include <unordered_map>
00005 #include <mutex>
00006 #include <list>
00007 #include <iostream>
00008 #include <map>
00009 #include <vector>
00010 #include <memory>
00011 #include <thread>
00012 #include <condition_variable>
00013 #include <atomic>
00014 #include <fstream>
00015 #include <filesystem>
00016 #include <queue>
00017 #ifdef __APPLE__
00018 #include <mach-o/dyld.h>
00019 #endif
00020
00021 class LRUCache {
00022 private:
00023     struct Node {
00024         std::string key;
00025         std::string value;
00026         Node* prev;
00027         Node* next;
00028         Node(const std::string& k, const std::string& v)
00029             : key(k), value(v), prev(nullptr), next(nullptr) {}
00030     };
00031
  
```

```

00032     size_t capacity_;
00033     std::unordered_map<std::string, Node*> cache_;
00034     Node* head_;
00035     Node* tail_;
00036     std::mutex mutex_;
00037
00038     void move_to_front(Node* node) {
00039         if (node == head_) return;
00040
00041         if (node == tail_) {
00042             tail_ = node->prev;
00043             tail_->next = nullptr;
00044         } else {
00045             node->prev->next = node->next;
00046             node->next->prev = node->prev;
00047         }
00048
00049         node->prev = nullptr;
00050         node->next = head_;
00051         head_->prev = node;
00052         head_ = node;
00053     }
00054
00055     void evict_lru() {
00056         if (tail_) {
00057             cache_.erase(tail_->key);
00058             Node* temp = tail_;
00059             tail_ = tail_->prev;
00060             if (tail_) tail_->next = nullptr;
00061             delete temp;
00062         }
00063     }
00064
00065 public:
00066     LRUcache(size_t capacity) : capacity_(capacity), head_(nullptr), tail_(nullptr) {}
00067
00068     ~LRUcache() {
00069         Node* current = head_;
00070         while (current) {
00071             Node* temp = current;
00072             current = current->next;
00073             delete temp;
00074         }
00075     }
00076
00077     size_t capacity() const { return capacity_; }
00078     size_t size() const { return cache_.size(); }
00079
00080     bool get(const std::string& key, std::string& value) {
00081         std::lock_guard<std::mutex> lock(mutex_);
00082         auto it = cache_.find(key);
00083         if (it != cache_.end()) {
00084             move_to_front(it->second);
00085             value = it->second->value;
00086             return true;
00087         }
00088         return false;
00089     }
00090
00091     bool put(const std::string& key, const std::string& value) {
00092         std::lock_guard<std::mutex> lock(mutex_);
00093         auto it = cache_.find(key);
00094         if (it != cache_.end()) {
00095             it->second->value = value;
00096             move_to_front(it->second);
00097             return true;
00098         }
00099
00100         if (cache_.size() >= capacity_) {
00101             evict_lru();
00102         }
00103
00104         Node* new_node = new Node(key, value);
00105         cache_[key] = new_node;
00106
00107         if (!head_) {
00108             head_ = tail_ = new_node;
00109         } else {
00110             new_node->next = head_;
00111             head_->prev = new_node;
00112             head_ = new_node;
00113         }
00114
00115         return true;
00116     }
00117
00118     bool remove(const std::string& key) {

```

```

00119     std::lock_guard<std::mutex> lock(mutex_);
00120     auto it = cache_.find(key);
00121     if (it != cache_.end()) {
00122         Node* node = it->second;
00123         if (node == head_) {
00124             head_ = node->next;
00125             if (head_) head_->prev = nullptr;
00126             else tail_ = nullptr;
00127         } else if (node == tail_) {
00128             tail_ = node->prev;
00129             if (tail_) tail_->next = nullptr;
00130         } else {
00131             node->prev->next = node->next;
00132             node->next->prev = node->prev;
00133         }
00134         cache_.erase(key);
00135         delete node;
00136         return true;
00137     }
00138     return false;
00139 }
00140 };
00141
00142 class DiskStorage {
00143 private:
00144     std::string data_file_;
00145     std::mutex mutex_;
00146     std::unordered_map<std::string, std::string> data_;
00147
00148     std::filesystem::path get_executable_path() {
00149 #ifdef __APPLE__
00150         char path[1024];
00151         uint32_t size = sizeof(path);
00152         if (_NSGetExecutablePath(path, &size) == 0) {
00153             return std::filesystem::path(path).parent_path();
00154         }
00155 #else
00156         char result[PATH_MAX];
00157         ssize_t count = readlink("/proc/self/exe", result, sizeof(result));
00158         if (count != -1) {
00159             return std::filesystem::path(result).parent_path();
00160         }
00161 #endif
00162     return std::filesystem::current_path();
00163 }
00164
00165 void ensure_directory_exists() {
00166     std::filesystem::path exe_path = get_executable_path();
00167     std::filesystem::path storage_dir = exe_path / "disk_storage";
00168     std::filesystem::create_directories(storage_dir);
00169 }
00170
00171 void load_data() {
00172     std::ifstream file(data_file_);
00173     if (!file.is_open()) return;
00174
00175     std::string line;
00176     while (std::getline(file, line)) {
00177         size_t pos = line.find('=');
00178         if (pos != std::string::npos) {
00179             std::string key = line.substr(0, pos);
00180             std::string value = line.substr(pos + 1);
00181             data_[key] = value;
00182         }
00183     }
00184     file.close();
00185 }
00186
00187 void save_data() {
00188     std::ofstream file(data_file_);
00189     if (!file.is_open()) return;
00190
00191     for (const auto& [key, value] : data_) {
00192         file << key << "=" << value << "\n";
00193     }
00194     file.close();
00195 }
00196
00197 public:
00198     DiskStorage() {
00199         try {
00200             ensure_directory_exists();
00201             std::filesystem::path exe_path = get_executable_path();
00202             data_file_ = (exe_path / "disk_storage" / "data.txt").string();
00203             load_data();
00204         } catch (const std::exception& e) {
00205             // If loading fails, start with empty data

```

```

00206         data_.clear();
00207     }
00208 }
00209
00210 ~DiskStorage() {
00211     save_data();
00212 }
00213
00214 bool get(const std::string& key, std::string& value) {
00215     std::lock_guard<std::mutex> lock(mutex_);
00216     auto it = data_.find(key);
00217     if (it != data_.end()) {
00218         value = it->second;
00219         return true;
00220     }
00221     return false;
00222 }
00223
00224 bool put(const std::string& key, const std::string& value) {
00225     std::lock_guard<std::mutex> lock(mutex_);
00226     data_[key] = value;
00227     save_data();
00228     return true;
00229 }
00230
00231 void remove(const std::string& key) {
00232     std::lock_guard<std::mutex> lock(mutex_);
00233     data_.erase(key);
00234     save_data();
00235 }
00236 };
00237
00238 class StorageEngine {
00239 public:
00240     StorageEngine(size_t cache_size = 1)
00241         : cache_(std::make_unique<LRUCache>(cache_size))
00242         , running_(false) {
00243     try {
00244         disk_storage_ = std::make_unique<DiskStorage>();
00245         running_ = true;
00246         write_thread_ = std::thread(&StorageEngine::async_write_worker, this);
00247     } catch (const std::exception& e) {
00248         running_ = false;
00249         throw;
00250     }
00251 }
00252
00253 ~StorageEngine() {
00254     if (running_) {
00255         running_ = false;
00256         write_cv_.notify_one();
00257         if (write_thread_.joinable()) {
00258             write_thread_.join();
00259         }
00260     }
00261 }
00262
00263 bool set(const std::string& key, const std::string& value) {
00264     return put(key, value);
00265 }
00266
00267 std::string get(const std::string& key) {
00268     std::string value;
00269     if (get(key, value)) {
00270         return value;
00271     }
00272     return "";
00273 }
00274
00275 bool get(const std::string& key, std::string& value) {
00276     // First check cache
00277     if (cache_->get(key, value)) {
00278         return true;
00279     }
00280
00281     // If not in cache, check disk storage
00282     if (disk_storage_->get(key, value)) {
00283         // Add to cache
00284         cache_->put(key, value);
00285         return true;
00286     }
00287     return false;
00288 }
00289
00290 bool del(const std::string& key) {
00291     try {
00292         // First check if key exists in either cache or disk

```

```

00293         std::string value;
00294         bool exists = false;
00295
00296         // Check cache first
00297         if (cache_->get(key, value)) {
00298             exists = true;
00299             cache_->remove(key);
00300         }
00301
00302         // Then check disk storage
00303         if (disk_storage_->get(key, value)) {
00304             exists = true;
00305             disk_storage_->remove(key);
00306         }
00307
00308         return exists;
00309     } catch (const std::exception& e) {
00310         return false;
00311     }
00312 }
00313
00314 void clear() {
00315     cache_ = std::make_unique<LRUCache>(cache_->capacity());
00316     disk_storage_ = std::make_unique<DiskStorage>();
00317 }
00318
00319 void force_flush() {
00320     std::lock_guard<std::mutex> lock(write_mutex_);
00321     while (!write_queue_.empty()) {
00322         auto [key, value] = write_queue_.front();
00323         write_queue_.pop();
00324         disk_storage_->put(key, value);
00325     }
00326 }
00327
00328 size_t size() const {
00329     return cache_->size();
00330 }
00331
00332 void sync() {
00333     force_flush();
00334 }
00335
00336 size_t pending_write_count() const {
00337     std::lock_guard<std::mutex> lock(const_cast<std::mutex&>(write_mutex_));
00338     return write_queue_.size();
00339 }
00340
00341 void stop_async_writer() {
00342     running_ = false;
00343     write_cv_.notify_one();
00344     if (write_thread_.joinable()) {
00345         write_thread_.join();
00346     }
00347 }
00348
00349 bool put(const std::string& key, const std::string& value) {
00350     // First try to put in cache
00351     if (cache_->put(key, value)) {
00352         // If successful, queue async write to disk
00353         {
00354             std::lock_guard<std::mutex> lock(write_mutex_);
00355             write_queue_.push({key, value});
00356         }
00357         write_cv_.notify_one();
00358         return true;
00359     }
00360
00361     // If cache is full, write directly to disk
00362     return disk_storage_->put(key, value);
00363 }
00364
00365 private:
00366     void async_write_worker() {
00367         while (running_) {
00368             std::unique_lock<std::mutex> lock(write_mutex_);
00369             write_cv_.wait(lock, [this] {
00370                 return !running_ || !write_queue_.empty();
00371             });
00372
00373             if (!running_) break;
00374
00375             auto [key, value] = write_queue_.front();
00376             write_queue_.pop();
00377             lock.unlock();
00378
00379             disk_storage_->put(key, value);
00380         }
00381     }
00382 }
```


```

00380         }
00381     }
00382
00383     std::unique_ptr<LRUCache> cache_;
00384     std::unique_ptr<DiskStorage> disk_storage_;
00385     std::queue<std::pair<std::string, std::string>> write_queue_;
00386     std::mutex write_mutex_;
00387     std::condition_variable write_cv_;
00388     std::thread write_thread_;
00389     std::atomic<bool> running_;
00390 };

```

## 5.9 part-b/src/main\_server.cpp File Reference

```
#include "network_server.h"
#include <iostream>
#include <signal.h>
Include dependency graph for main_server.cpp:
```



### Functions

- void [signal\\_handler](#) (int signum)
- int [main](#) ()

### Variables

- [Server \\* g\\_server = nullptr](#)

## 5.9.1 Function Documentation

### 5.9.1.1 main()

```
int main ()
```

Here is the call graph for this function:



### 5.9.1.2 signal\_handler()

```
void signal_handler (
    int signum)
```

Here is the caller graph for this function:



## 5.9.2 Variable Documentation


### 5.9.2.1 g\_server

```
Server* g_server = nullptr
```

## 5.10 part-b/src/network\_client.cpp File Reference

```
#include <iostream>
#include <string>
#include <cstring>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <unistd.h>
#include <stdexcept>
#include <sstream>
```

Include dependency graph for network\_client.cpp:



## Classes

- class [NetworkClient](#)

## Functions

- int [main](#) (int argc, char \*argv[ ])

### 5.10.1 Function Documentation

### 5.10.1.1 main()


```
int main (
    int argc,
    char * argv[ ]) {
```

Here is the call graph for this function:



## 5.11 part-b/src/network\_server.cpp File Reference

```
#include "network_server.h"
#include <iostream>
#include <sstream>
#include <vector>
#include <cstring>
#include <fcntl.h>
#include <netinet/tcp.h>
#include <sys/socket.h>
#include <sys/event.h>
#include <unistd.h>
#include <arpa/inet.h>
```



## Macros

- #define TCP\_NODELAY 1


### 5.11.1 Macro Definition Documentation

#### 5.11.1.1 TCP NODELAY


```
#define TCP_NODELAY 1
```

## 5.12 part-b/src/network\_server.h File Reference

```
#include "storage_engine.h"
#include <memory>
#include <string>
#include <unordered_map>
Include dependency graph for network_server.h:
```



This graph shows which files directly or indirectly include this file:



### Classes

- class [Server](#)

## 5.13 network\_server.h

[Go to the documentation of this file.](#)

```
00001 #pragma once
00002
00003 #include "storage_engine.h"
00004 #include <memory>
00005 #include <string>
00006 #include <unordered_map>
00007
00008 class Server {
00009 public:
00010     Server();
00011     ~Server();
00012
00013     void run();
00014     void stop();
00015 }
```

```
00016 private:
00017     static constexpr int PORT = 9001;
00018     static constexpr int LISTEN_BACKLOG = 128;
00019     static constexpr int MAX_EVENTS = 1024;
00020
00021     int server_fd = -1;
00022     int kq = -1;
00023     bool should_stop = false;
00024     std::unique_ptr<StorageEngine> storage;
00025     std::unordered_map<int, std::string> client_buffers;
00026
00027     void setup_server();
00028     void setup_kqueue();
00029     void set_nonblocking(int fd);
00030     void handle_new_connection();
00031     void handle_client_data(int client_fd);
00032     std::string process_command(const std::string& command);
00033     std::string encode_resp(const std::string& response);
00034 }
```



# Index

~DiskStorage  
    DiskStorage, 15  
~LRUCache  
    LRUCache, 19  
~NetworkClient  
    NetworkClient, 23  
~Server  
    Server, 26  
~StorageEngine  
    StorageEngine, 33

access\_order  
    StorageEngine, 42  
async\_write\_worker  
    StorageEngine, 34

BATCH\_SIZE  
    StorageEngine, 42

BLINK DB Documentation, 1

cache  
    LRUCache, 22  
    StorageEngine, 42

capacity  
    LRUCache, 20

capacity\_  
    LRUCache, 22

clear  
    StorageEngine, 34

client\_buffers  
    Server, 30

data\_  
    DiskStorage, 18  
    StorageEngine, 42

DATA\_FILE  
    StorageEngine, 42

data\_file\_  
    DiskStorage, 18

del  
    StorageEngine, 34, 35

DISK\_DIR  
    StorageEngine, 42

disk\_index  
    StorageEngine, 43

disk\_storage\_  
    StorageEngine, 43

DiskStorage, 14  
    ~DiskStorage, 15  
    data\_, 18

    data\_file\_, 18  
    DiskStorage, 15  
    ensure\_directory\_exists, 15  
    get, 16  
    get\_executable\_path, 16  
    load\_data, 16  
    mutex\_, 18  
    put, 16  
    remove, 17  
    save\_data, 17

    docs/mainpage.md, 45

    encode\_resp  
        Server, 27

    ensure\_directory\_exists  
        DiskStorage, 15

    evict\_lru  
        LRUCache, 20

    flush\_write\_buffer  
        StorageEngine, 35

    force\_flush  
        StorageEngine, 36

    g\_server  
        main\_server.cpp, 56

    get  
        DiskStorage, 16  
        LRUCache, 20  
        StorageEngine, 36, 37

    get\_executable\_path  
        DiskStorage, 16

    handle\_client\_data  
        Server, 27

    handle\_new\_connection  
        Server, 27

    head\_  
        LRUCache, 22

    INDEX\_FILE  
        StorageEngine, 43

    key  
        LRUCache::Node, 25  
        StorageEngine::BatchEntry, 13

    kq  
        Server, 30

    LISTEN\_BACKLOG  
        Server, 31

load\_data  
     DiskStorage, 16  
 load\_disk\_index  
     StorageEngine, 37  
 LRUcache, 18  
     ~LRUcache, 19  
     cache\_, 22  
     capacity, 20  
     capacity\_, 22  
     evict\_lru, 20  
     get, 20  
     head\_, 22  
     LRUcache, 19  
     move\_to\_front, 20  
     mutex\_, 22  
     put, 21  
     remove, 21  
     size, 21  
     tail\_, 22  
 LRUcache::Node, 24  
     key, 25  
     next, 25  
     Node, 25  
     prev, 25  
     value, 25  
  
 main  
     main\_server.cpp, 55  
     network\_client.cpp, 57  
     repl.cpp, 46  
 main\_server.cpp  
     g\_server, 56  
     main, 55  
     signal\_handler, 55  
 MAX\_CACHE\_SIZE  
     StorageEngine, 43  
 MAX\_EVENTS  
     Server, 31  
 MAX\_KEY\_SIZE  
     StorageEngine, 43  
 MAX\_VALUE\_SIZE  
     StorageEngine, 43  
 move\_to\_front  
     LRUcache, 20  
 mutex\_  
     DiskStorage, 18  
     LRUcache, 22  
     StorageEngine, 43  
  
 network\_client.cpp  
     main, 57  
 network\_server.cpp  
     TCP\_NODELAY, 57  
 NetworkClient, 22  
     ~NetworkClient, 23  
     NetworkClient, 23  
     parse\_response, 23  
     send\_command, 23  
     server\_addr, 24  
  
     sock, 24  
 next  
     LRUCache::Node, 25  
 Node  
     LRUCache::Node, 25  
  
 offset  
     StorageEngine::DiskEntry, 14  
  
 parse\_response  
     NetworkClient, 23  
 part-a/src/repl.cpp, 45  
 part-a/src/storage\_engine.cpp, 46  
 part-a/src/storage\_engine.h, 47, 48  
 part-b/src/main\_server.cpp, 55  
 part-b/src/network\_client.cpp, 56  
 part-b/src/network\_server.cpp, 57  
 part-b/src/network\_server.h, 58  
 part-b/src/storage\_engine.cpp, 47  
 part-b/src/storage\_engine.h, 49, 50  
 pending\_write\_count  
     StorageEngine, 38  
 pending\_writes  
     StorageEngine, 43  
 PORT  
     Server, 31  
 prev  
     LRUCache::Node, 25  
 printUsage  
     repl.cpp, 46  
 process\_command  
     Server, 28  
 put  
     DiskStorage, 16  
     LRUcache, 21  
     StorageEngine, 38  
  
 remove  
     DiskStorage, 17  
     LRUcache, 21  
 remove\_from\_disk\_index  
     StorageEngine, 38  
 repl.cpp  
     main, 46  
     printUsage, 46  
 run  
     Server, 28  
 running\_  
     StorageEngine, 43  
  
 save\_data  
     DiskStorage, 17  
 save\_disk\_index  
     StorageEngine, 39  
 send\_command  
     NetworkClient, 23  
 Server, 25  
     ~Server, 26  
     client\_buffers, 30

encode\_resp, 27  
handle\_client\_data, 27  
handle\_new\_connection, 27  
kq, 30  
LISTEN\_BACKLOG, 31  
MAX\_EVENTS, 31  
PORT, 31  
process\_command, 28  
run, 28  
Server, 26  
server\_fd, 31  
set\_nonblocking, 29  
setup\_kqueue, 29  
setup\_server, 29  
should\_stop, 31  
stop, 30  
storage, 31  
server\_addr  
    NetworkClient, 24  
server\_fd  
    Server, 31  
set  
    StorageEngine, 39, 40  
set\_nonblocking  
    Server, 29  
setup\_kqueue  
    Server, 29  
setup\_server  
    Server, 29  
should\_stop  
    Server, 31  
signal\_handler  
    main\_server.cpp, 55  
size  
    LRUCache, 21  
    StorageEngine, 40, 41  
    StorageEngine::DiskEntry, 14  
sock  
    NetworkClient, 24  
stop  
    Server, 30  
stop\_async\_writer  
    StorageEngine, 41  
storage  
    Server, 31  
StorageEngine, 31  
    ~StorageEngine, 33  
    access\_order, 42  
    async\_write\_worker, 34  
    BATCH\_SIZE, 42  
    cache\_, 42  
    clear, 34  
    data\_, 42  
    DATA\_FILE, 42  
    del, 34, 35  
    DISK\_DIR, 42  
    disk\_index, 43  
    disk\_storage\_, 43  
    flush\_write\_buffer, 35  
    force\_flush, 36  
    get, 36, 37  
    INDEX\_FILE, 43  
    load\_disk\_index, 37  
    MAX\_CACHE\_SIZE, 43  
    MAX\_KEY\_SIZE, 43  
    MAX\_VALUE\_SIZE, 43  
    mutex\_, 43  
    pending\_write\_count, 38  
    pending\_writes, 43  
    put, 38  
    remove\_from\_disk\_index, 38  
    running\_, 43  
    save\_disk\_index, 39  
    set, 39, 40  
    size, 40, 41  
    stop\_async\_writer, 41  
    StorageEngine, 33  
    sync, 41  
    update\_disk\_index, 41  
    write\_buffer, 43  
    write\_cv\_, 44  
    write\_mutex\_, 44  
    write\_queue\_, 44  
    write\_thread\_, 44  
    StorageEngine::BatchEntry, 13  
        key, 13  
        value, 13  
    StorageEngine::DiskEntry, 13  
        offset, 14  
        size, 14  
    sync  
        StorageEngine, 41  
    tail\_  
        LRUCache, 22  
TCP\_NODELAY  
    network\_server.cpp, 57  
update\_disk\_index  
    StorageEngine, 41  
value  
    LRUCache::Node, 25  
    StorageEngine::BatchEntry, 13  
write\_buffer  
    StorageEngine, 43  
write\_cv\_  
    StorageEngine, 44  
write\_mutex\_  
    StorageEngine, 44  
write\_queue\_  
    StorageEngine, 44  
write\_thread\_  
    StorageEngine, 44